


CS 188

Scalable Internet Services

John Rothfels
October 29, 2019



For Today



Motivation

After today's lecture you will understand how NoSQL can be used to build scalable internet services.

NoSQL data stores won't be part of your project, but after this lecture you should understand when you could use them.



Motivation



As our application has experienced greater and greater popularity, the data layer has proven difficult to scale horizontally.

Without a scaling path for our data layer, it will be a bottleneck limiting our application.




Motivation

Relational Databases are great tools for our data layer, but we can't simply spread load across multiple RDBMSes.

We will look at a few techniques for scaling RDBMSes next time:

- Sharding
- Service Oriented Architectures
- Distinguishing Reads from Writes




Motivation

What if these techniques aren't sufficient for our target application?

- There's no good way to shard our application?
- We've already broken our application out via SOA and still have load hotspots?
- We're already using read-slaves and it's not enough?

When relational databases fail to scale to our needs, we need to turn to non-relational solutions.



NoSQL



Non-relational databases are sometimes called NoSQL databases.

This is an umbrella term for many types of databases

- Key-value stores
- Column-oriented data stores
- Document-oriented stores
- Graph databases



NoSQL




Most NoSQL solutions are good at horizontal scaling.

- You can easily add hardware to the database to increase throughput.

In exchange for better horizontal scaling, these databases provide the application layer fewer guarantees



NoSQL




Let's say we want a database to span multiple machines.

We can update on both nodes, and the databases keep each other in synch.



NoSQL



Lets say:

- We have two clients sending writes.
- There's a uniqueness constraint on email.


If both try to write the same email address to different rows, and the databases can communicate, they can resolve this in some manner.

- Ex: Allow one, fail the other.




NoSQL

How do we handle a network partition?



NoSQL




How do we handle a network partition?

- If the databases can't communicate, they don't know if this update violates database consistency.
- Allow the write and hope for the best?
 - And fix it later if needed?
- Not accept such writes during a partition?



NoSQL




If we allow the write, our database is not consistent.

If we don't allow the write, our database is unavailable.

If we ignore this scenario, we can't tolerate network partitions.



NoSQL




It is situations like these that motivated the CAP theorem.

- Eric Brewer, late 90s.
- “Consistency, Availability, Partition Tolerance: choose any two.”



NoSQL




AP systems:

- Always up
- Can handle network partitions
- Not always consistent.

In the earlier example, an AP solution would accept the writes.



NoSQL




CP systems:

- Always consistent.
- Handle network partitions.
- Sometimes will be unavailable to clients

In the earlier example, a CP solution would not allow any writes



NoSQL




CA systems:

- Always up.
- Always consistent.
- Assume no network partitions.

A CA solution would never get into the earlier scenario because it wouldn't be deployed where partitions could happen.



NoSQL



Assuming no partitions is very limiting

- For HA and latency reasons, we'd like to have multi-site
- Even within a datacenter, we could have a partition

As a result, trade-offs tend to be more C vs A



ACID vs. BASE

The **BASE** acronym was created to describe these NoSQL solutions that make tradeoffs between Availability and Consistency

ACID

Atomicity

Consistency

Isolation

Durability

BASE

Basically

Available

Soft State

Eventually Consistent



Consistency

Consistency comes in many forms:

- Strong Consistency
 - After update, everyone sees new value
- Eventual Consistency
 - Eventually the system will converge on the new value
 - Read-your-writes Consistency
 - You immediately see any data you have written
 - Causal Consistency
 - You see your own writes, and anyone you communicate with sees your writes
 - Session Consistency
 - Within a session, you see your own writes.



NWR



Web Server Process
Operating System



N, W & R are a useful shorthand for describing the read/write strategy of a data store.



NWR




N refers to the number of separate nodes that each retain a copy of the data.

In this example, N is 5.



NWR




W refers to the number of nodes that we persist to before considering a write written.

In this example, W is 3.



NWR




R refers to the number of nodes that we consult when reading.

In this example, R is 2.




NWR



If $W + R \leq N$, then you can't be sure that a read has seen all previous writes



NWR




If $W + R \leq N$, then you can't be sure that a read has seen all previous writes.

$$3 + 2 \leq 5.$$



NWR



If $W + R > N$, then you can be sure that a read has seen all previous writes.

Any two size-3 subsets of 5 servers must have overlap

$$3 + 3 > 5.$$



NWR

For strong consistency, many combinations can work.

- $R=N, W=1$: Write to any one server, consult all server on reads. Use the newest value.
- $R=1, W=N$: Write to all servers, consult any server on reads.
- $R = N/2 + 1, W = N/2 + 1$: Write to a quorum, read from a quorum.



NWR

For weaker notions of consistency, we choose $W + R \leq N$.

Exactly which type of consistency we see will depend on “session stickiness”

- I can write to multiple nodes, but if User A's reads and writes to the same node, we can more easily implement “see your own writes” consistency.
- Similar use of stickiness to achieve Session Consistency.



NoSQL

There are different types of NoSQL stores

- Document-oriented stores
 - We will look at MongoDB
- Key-value stores
 - We will look at Redis
- Column-oriented data stores
 - We will look at Cassandra
- Graph databases
 - We won't be looking at these today
 - Specialized data stores, not always horizontally scalable.



MongoDB



MongoDB is a Document-oriented data store.

- Stores “Documents” that are nested hash-like structures.
- These Documents are stored in “Collections” (similar to a table in RDBMS).
- Has no fixed Schema.
- Docs can have references to other docs



Collection



MongoDB

Collection
↓
`db.users.insert({`
Document
↓
 `name: "sue",`
 `age: 26,`
 `status: "A",`
 `groups: ["news", "sports"]`
`})`

Query language is not SQL

Document

```
{
  name: "sue",
  age: 26,
  status: "A",
  groups: [ "news", "sports" ]
}
```

insert →

Collection

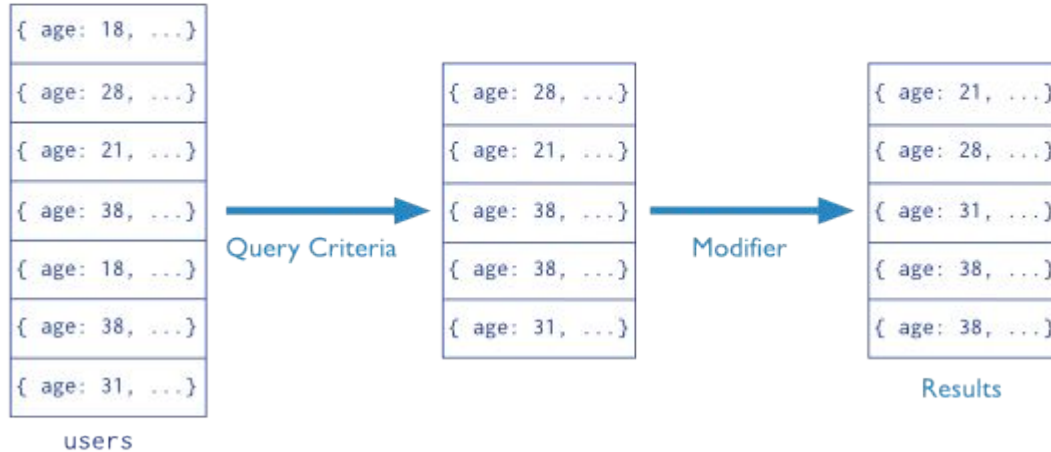
{ name: "al", age: 18, ... }
{ name: "lee", age: 28, ... }
{ name: "jan", age: 21, ... }
{ name: "kai", age: 38, ... }
{ name: "sam", age: 18, ... }
{ name: "mel", age: 38, ... }
{ name: "ryan", age: 31, ... }
{ name: "sue", age: 26, ... }

users



MongoDB

Collection Query Criteria Modifier
`db.users.find({ age: { $gt: 18 } }).sort({age: 1 })`



Query language is not SQL



MongoDB

Documents are stored in JSONB

- Binary version of JSON
- Can nest other JSON documents



No notion of transactions

- Unit of atomicity is the Document.
- Inserting multiple documents can fail individually



MongoDB

No notion of Joins.

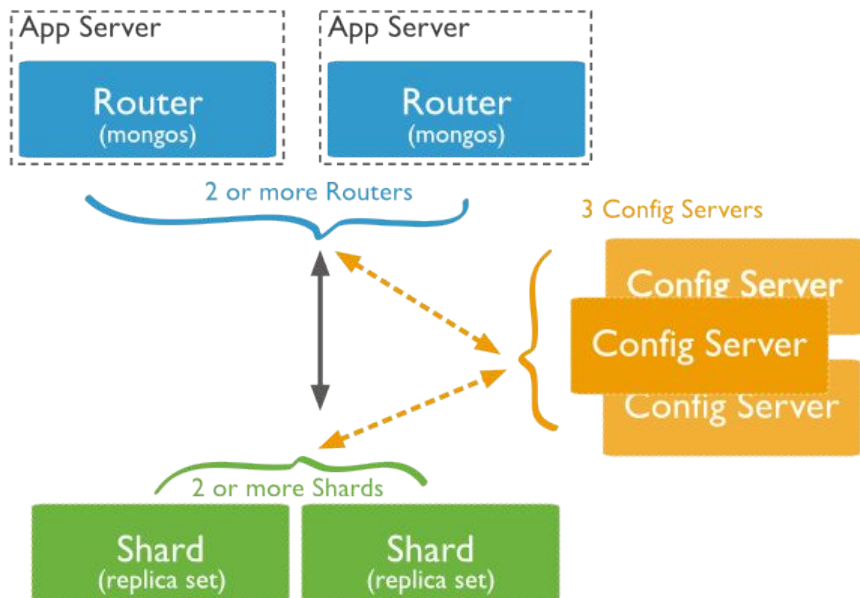
- If you want to do computation based on relations between documents, read them into memory and do them at the application layer.



Can have secondary indexes based on document values



MongoDB




Collections can be sharded

- Each shard can have a replica set
- Config Servers manage the mapping between shards and data.
- Mongo routes queries to the appropriate shard



MongoDB



Replica sets use asynchronous replication

You can configure your driver to read from the primary only, or to read from read-replicas


- Reading from Primary:
($R=1$, $W=1$, $N=1$)
- Reading from replicas:
($R=1$, $W=1$, $N=3$)



Redis

Redis is a key-value data store.

- Also called a data structure store
- Supports many data structures
 - Lists
 - Sorted Sets
 - Hashes
 - Bitmaps



redis

Primarily keeps data structures
in memory


- Persistence to disk is
optional



Redis

Redis interface

- Each data type allows similar mechanisms to what you would do in memory
- Access hashes by key
- Access lists by index
- Sorted sets can return top-K
- Push/pop on lists




redis



Redis

Redis sort of has transactions

- Redis operations are simple
- You can batch up a series of commands into a transaction
- When a command fails, the previous do not roll back




redis



Redis

Two options for disk persistence


- RDF: Redis Database File
 - Forks process and saves a dump
- AOF: Append Only File
 - Saves updates to a log
 - Log is replayed upon start



redis



Redis




Redis cluster supports sharding

Single master for writes, replicas for failover.



Redis




Cluster can handle all reads if up to two nodes are down.

It's possible to read from slaves, but default sends all read and write operations to master.



Redis

Cluster can also dynamically rebalance after adding hardware.



Cassandra

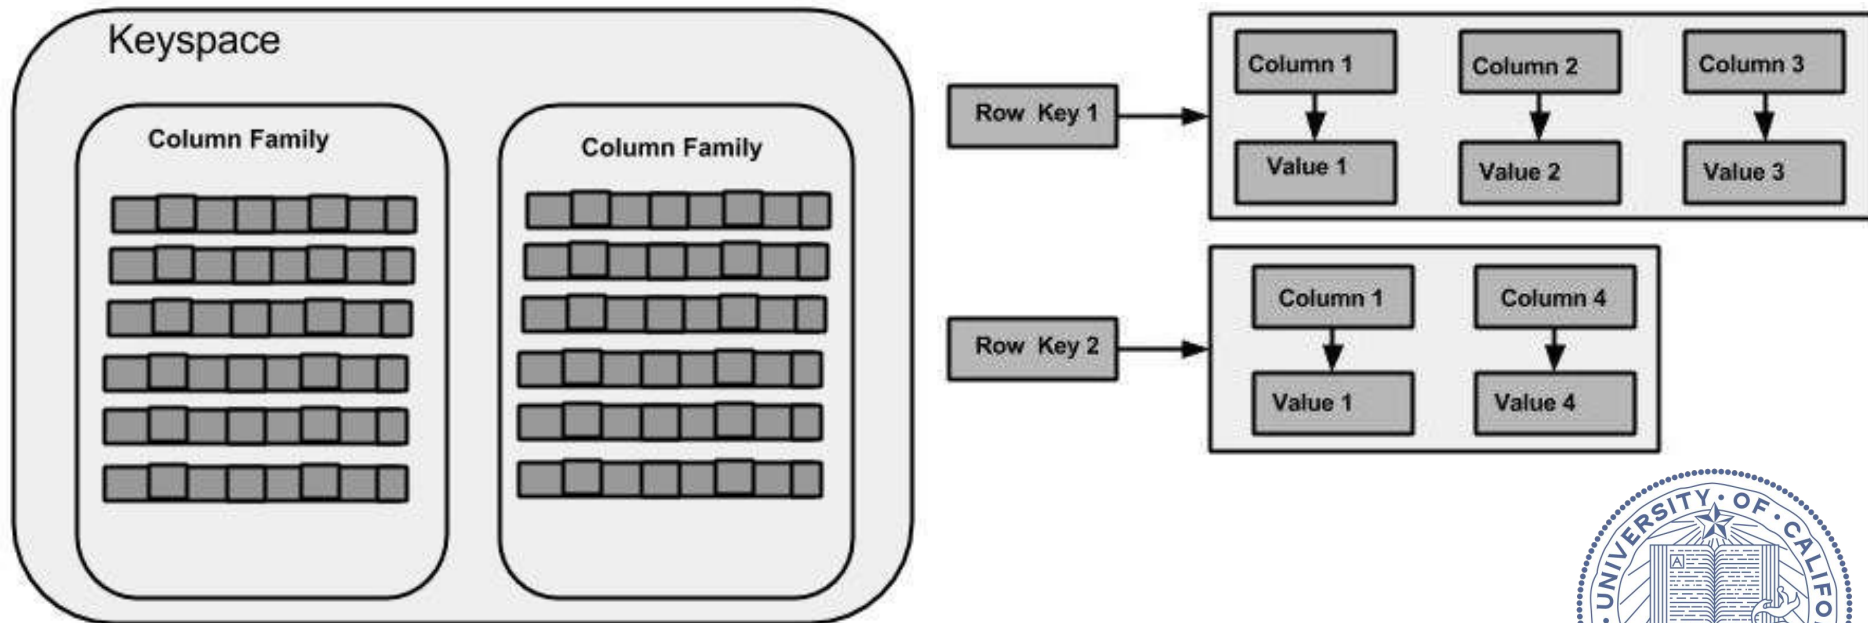
Most similar to a RDBMS

Cassandra has table-like structures called ColumnFamilies

- ColumnFamilies have many rows
- Rows are like a big hash, with many keys and values
- Rows are heterogeneous and can be schemaless
- Rows can be very long
 - Have many keys and values



Cassandra



Cassandra

Interface is called CQL, similar to SQL

```
SELECT * WHERE KEY = 11194251 AND  
startdate = '2011-10-08-0500';
```

Features are very limited

- Most queries are key-value
- Secondary indices are allowed
- Sorting is very limited




Cassandra

No transactions

- Atomic batches exist
- No isolation from other batches

No Joins.

- Do this at the application layer



cassandra



Cassandra

- Cassandra is a masterless system
- Distributed and highly available
- Data is automatically split across nodes
- Reads are eventually consistent
 - But can be made strictly consistent (per statement)



Cassandra

- Consistency per statement

SELECT * WHERE KEY = 11194251...

CONSISTENCY LEVEL ONE (R=1)

CONSISTENCY LEVEL ALL (R=N)

CONSISTENCY LEVEL QUORUM (R=N/2+1)

UPDATE ... WHERE KEY = 11194251...

CONSISTENCY LEVEL ONE (W=1)

CONSISTENCY LEVEL ALL (W=N)

CONSISTENCY LEVEL QUORUM (W=N/2+1)



Motivation

After today's lecture you will understand how NoSQL can be used to build into scalable internet services.

NoSQL data stores won't likely be part of your project, but after this lecture you should understand when you could look to them in practice.



For Next Time...

Next lecture we will discuss how to horizontally scale RDBMS.

You should be building lots of features and functionality in your projects now.

Piazza thread for AWS troubleshooting.

